

Modifying CFS Scheduler

Samiksha Modi
Roll No: 2019331

Description of Code and How it Works

I defined my system call in `rtnice.c` present in the directory `rtnice` in the extracted folder of the kernel v5.9.1

The function is defined in the following way : `SYSCALL_DEFINEX()`, where X is the number of parameters that are passed to the system call. For me it was `SYSCALL_DEFINE2(rtnice, int, mypid, int, mytime)`. Here $x=2$ since I have 2 parameters ie: `pid (mypid)` and `soft realtime guarantee (mytime)`. The first parameter in the definition is compulsory and is the name of the system call implemented. I then multiply `mytime` with 10^{11} to get a significantly large value for `soft_rtnice`. If `mytime` is a negative value, I return the error no and handle the error.

To get struct `pid` I use `find_get_pid()`. If it returns `NULL` then no process with the given input `pid` exists, so I return the error no and handle the error.

If it exists, then I use `pid_task()` to get the `task_struct` of the given `pid`. The `task_struct` contains `sched_entity (se)` of the process in which I assign the soft realtime guarantee called `new`.

To add the data variable `soft_rtnice` in the `task_struct`, I created a new data variable in the file `sched.h`, and initialized it to 0 in `core.c`.

To guarantee that the process with the given `pid` runs for the given time, we need to first modify the scheduler in such a way that the scheduler always picks a task with lower `soft_rtnice`, and if two processes have the same `soft_rtnice`, then the priority should be given to the task with lower `vruntime`. I did this in `entity_before()` present in `fair.c`

I also deduct the amount of time the program has run from the `soft_rtnice`. For this I decrement the `delta_exec` value from the `soft_rtnice`, if the value of `delta_exec` is greater than the amount of `soft_rtnice`, then I set the `soft_rtnice` to 0.

User Inputs

- PID (datatype- int)
- Soft realtime guarantee (datatype-int)

The inputs have been hardcoded into `test.c`

The system can be tested by changing the arguments we pass in `syscall`.

```

34     printf("Time 2: %lf\n",time);
35     exit(EXIT_SUCCESS);
36 }
37 else
38 {
39     long ans=syscall(440,getpid(),100);
40     //long ans=syscall(440,-2,100);
41     //long ans=syscall(440,getpid(),-4);
42 }

```

The arguments in syscall are syscall number (440), pid, and soft real-time guarantee.

Expected output

I am printing the execution time taken by both the tasks. I fork() my process and in one process I print the execution time with soft realtime guarantees. And in the other process I print the execution time without the soft realtime guarantees. Values of both the execution time will be different. Incase of an invalid input, it will show the appropriate error.

Errors Handled

1. Invalid pid input

It will show Error: 2. It means that no process with the given pid exists.

```

sam@sam-VirtualBox:~/Desktop$ make run
gcc test.c
./a.out
Error: 2
sam@sam-VirtualBox:~/Desktop$ Time 2: 3.772034

```

2. Invalid time input

It will show Error: 3. If the user inputs a negative value for the soft realtime guarantee, it gives an error.

```

sam@sam-VirtualBox:~/Desktop$ make run
gcc test.c
./a.out
Error: 3
sam@sam-VirtualBox:~/Desktop$ Time 2: 4.465869

```

Sample 1 (Invalid pid input)

```
34     printf("Time 2: %lf\n",time);
35     exit(EXIT_SUCCESS);
36 }
37 else
38 {
39     //Long ans=syscall(440,getpid(),100);
40     Long ans=syscall(440,-2,100);
41     //Long ans=syscall(440,getpid(),-4);
42 }
```

```
sam@sam-VirtualBox:~/Desktop$ make run
gcc test.c
./a.out
Error: 2
sam@sam-VirtualBox:~/Desktop$ Time 2: 3.772034
```

Sample 2 (Invalid time input)

```
34     printf("Time 2: %lf\n",time);
35     exit(EXIT_SUCCESS);
36 }
37 else
38 {
39     //Long ans=syscall(440,getpid(),100);
40     //Long ans=syscall(440,-2,100);
41     Long ans=syscall(440,getpid(),-4);
42 }
```

```
sam@sam-VirtualBox:~/Desktop$ make run
gcc test.c
./a.out
Error: 3
sam@sam-VirtualBox:~/Desktop$ Time 2: 4.465869
```

Sample 3 (Correct user inputs)

```
34     printf("Time 2: %lf\n",time);
35     exit(EXIT_SUCCESS);
36 }
37 else
38 {
39     long ans=syscall(440,getpid(),100);
40     //long ans=syscall(440,-2,100);
41     //long ans=syscall(440,getpid(),-4);
42 }
```

```
sam@sam-VirtualBox:~/Desktop$ make run
gcc test.c
./a.out
Time 1: 3.796540
Time 2: 7.432829
sam@sam-VirtualBox:~/Desktop$
```